

Code ▾

# Regression

Linear Regression is a method used to find similarities between a dependent variable and 1 or more independent variable. The goal of linear regression model is to predict the dependent value by using the independent values. Some strengths of linear regression may include that it is widely used and it can be used to understand relationships between objects. One weakness may include that it assumes a linear relationship but it may not be the case all the time.

Hide

```
set.seed(3)
data <- read.csv(file = 'desktop/creditcard.csv')
dim(data)
```

```
[1] 284807    31
```

Hide

```
head(data)
```

T...		V1	V2	V3	V4	V5	V6	V7
<dbl>		<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0	-1.3598071	-0.07278117	2.5363467	1.3781552	-0.33832077	0.46238778	0.23959855
2	0	1.1918571	0.26615071	0.1664801	0.4481541	0.06001765	-0.08236081	-0.07880298
3	1	-1.3583541	-1.34016307	1.7732093	0.3797796	-0.50319813	1.80049938	0.79146096
4	1	-0.9662717	-0.18522601	1.7929933	-0.8632913	-0.01030888	1.24720317	0.23760894
5	2	-1.1582331	0.87773675	1.5487178	0.4030339	-0.40719338	0.09592146	0.59294075
6	2	-0.4259659	0.96052304	1.1411093	-0.1682521	0.42098688	-0.02972755	0.47620095

6 rows | 1-9 of 31 columns

Add a new chunk by clicking the *Insert Chunk* button on the toolbar or by pressing *Cmd+Option+I*.

Hide

```
set.seed(3)
i<-sample(1:nrow(data), nrow(data)*0.8, replace = FALSE)
train<-data[i,]
test<-data[-i,]
```

When you save the notebook, an HTML file containing the code and output will be saved alongside it (click the *Preview* button or press *Cmd+Shift+K* to preview the HTML file).

Hide

```
dim(train)
```

```
[1] 227845      31
```

Hide

```
head(train)
```

Time		V1	V2	V3	V4	V5	V6	
<dbl>		<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	
218087	141166	1.7933193	0.01984035	-1.9296190	1.34779239	0.5610834	-1.3402134	1.0
236218	148714	-0.8220090	0.72748442	-0.8316150	0.98524199	2.2795032	-0.5769322	0.9
229706	146020	2.0397723	0.61324571	-2.3616822	0.76673928	0.5995394	-1.7172058	0.3
83869	60061	-0.5231397	0.58988326	3.2996937	3.05804626	-0.4433193	1.5899776	-0.2
197613	132097	1.8249885	-0.68401799	-0.7133819	0.09200484	-0.3199603	0.1725305	-0.4
268264	163135	-0.2646561	1.00158569	-1.1063368	-1.25530667	1.5773387	-0.5076532	1.6

6 rows | 1-9 of 31 columns

Hide

```
names(train)
```

```
[1] "Time"      "v1"      "v2"      "v3"      "v4"      "v5"      "v6"
[8] "v7"      "v8"      "v9"      "v10"     "v11"     "v12"     "v13"
[15] "v14"     "v15"     "v16"     "v17"     "v18"     "v19"     "v20"
[22] "v21"     "v22"     "v23"     "v24"     "v25"     "v26"     "v27"
[29] "v28"     "Amount"  "Class"
```

Hide

```
colSums(is.na(train))
```

Time	V1	V2	V3	V4	V5	V6	V7	V8	V9
0	0	0	0	0	0	0	0	0	0
V10	V11	V12	V13	V14	V15	V16	V17	V18	V19
0	0	0	0	0	0	0	0	0	0
V20	V21	V22	V23	V24	V25	V26	V27	V28	Amount
0	0	0	0	0	0	0	0	0	0
Class									
0									

Hide

```
str(train)
```

```
'data.frame':  227845 obs. of  31 variables:
 $ Time   : num  141166 148714 146020 60061 132097 ...
 $ V1     : num  1.793 -0.822 2.04 -0.523 1.825 ...
 $ V2     : num  0.0198 0.7275 0.6132 0.5899 -0.684 ...
 $ V3     : num  -1.93 -0.832 -2.362 3.3 -0.713 ...
 $ V4     : num  1.348 0.985 0.767 3.058 0.092 ...
 $ V5     : num  0.561 2.28 0.6 -0.443 -0.32 ...
 $ V6     : num  -1.34 -0.577 -1.717 1.59 0.173 ...
 $ V7     : num  1.081 0.938 0.392 -0.201 -0.491 ...
 $ V8     : num  -0.619 -0.118 -0.411 0.362 0.133 ...
 $ V9     : num  -0.356 -0.25 0.439 -0.294 0.779 ...
 $ V10    : num  0.267 0.222 -1.221 0.575 0.123 ...
 $ V11    : num  -0.4948 0.4926 0.3239 0.0458 0.5658 ...
 $ V12    : num  0.8447 -0.9152 0.0279 0.3468 0.8999 ...
 $ V13    : num  1.1326 -2.2665 0.4242 -0.3725 0.0711 ...
 $ V14    : num  0.604 -0.742 -2.618 -0.885 0.137 ...
 $ V15    : num  0.1103 -0.487 1.0684 -1.4965 -0.0971 ...
 $ V16    : num  -0.512 -0.397 0.692 -0.442 0.758 ...
 $ V17    : num  -0.448 0.94 1.8 0.255 -0.902 ...
 $ V18    : num  -0.635 1.131 1.126 0.572 0.174 ...
 $ V19    : num  -0.509 0.996 -0.773 1.813 0.379 ...
 $ V20    : num  0.0696 0.1527 -0.1313 0.272 0.041 ...
 $ V21    : num  0.2258 -0.0806 0.1085 -0.0874 -0.188 ...
 $ V22    : num  0.519 0.237 0.546 0.209 -0.718 ...
 $ V23    : num  -0.1445 -0.4062 -0.0462 -0.2926 0.3409 ...
 $ V24    : num  0.0275 0.1415 -0.2154 0.0288 0.3405 ...
 $ V25    : num  0.4489 0.0203 0.2581 0.0884 -0.6436 ...
 $ V26    : num  -0.5 -0.413 -0.099 0.294 0.193 ...
 $ V27    : num  -0.0471 0.4572 0.0262 -0.0263 -0.066 ...
 $ V28    : num  -0.03701 0.09948 0.00825 -0.13511 -0.03495 ...
 $ Amount: num  144 21 1 29 98 ...
 $ Class  : int  0 0 0 0 0 0 0 0 0 0 ...
```

Hide

```
install.packages("vioplot")
```

```
trying URL 'https://cran.rstudio.com/bin/macosx/contrib/4.2/vioplot_0.4.0.tgz'
Content type 'application/x-gzip' length 1385815 bytes (1.3 MB)
=====
downloaded 1.3 MB
```

```
The downloaded binary packages are in
  /var/folders/fx/fht6tvb95xldy2488rl0ctj00000gn/T//RtmpNP8aNE/downloaded_packages
```


Hide

```
library("vioplot")
```


```
Loading required package: sm  
Package 'sm', version 2.2-5.7: type help(sm) for summary information  
Loading required package: zoo  
  
Attaching package: 'zoo'  
  
The following objects are masked from 'package:base':  
  
    as.Date, as.Date.numeric
```

[Hide](#)

```
vioplot(train$Time, xlab = "Time", ylab = "transaction")
```

[Hide](#)

```
boxplot(train$Time, train$Amount, xlab="Time", ylab = "Amount")
```


[Hide](#)

```
lm1<-lm(Time~Amount, data = train)
summary(lm1)
```

Call:

```
lm(formula = Time ~ Amount, data = train)
```

Residuals:

Min	1Q	Median	3Q	Max
-95058	-40574	-10074	44482	121125

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	95062.9008	105.3560	902.302	< 2e-16 ***
Amount	-1.9458	0.3929	-4.952	7.35e-07 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 47480 on 227843 degrees of freedom

Multiple R-squared: 0.0001076, Adjusted R-squared: 0.0001032

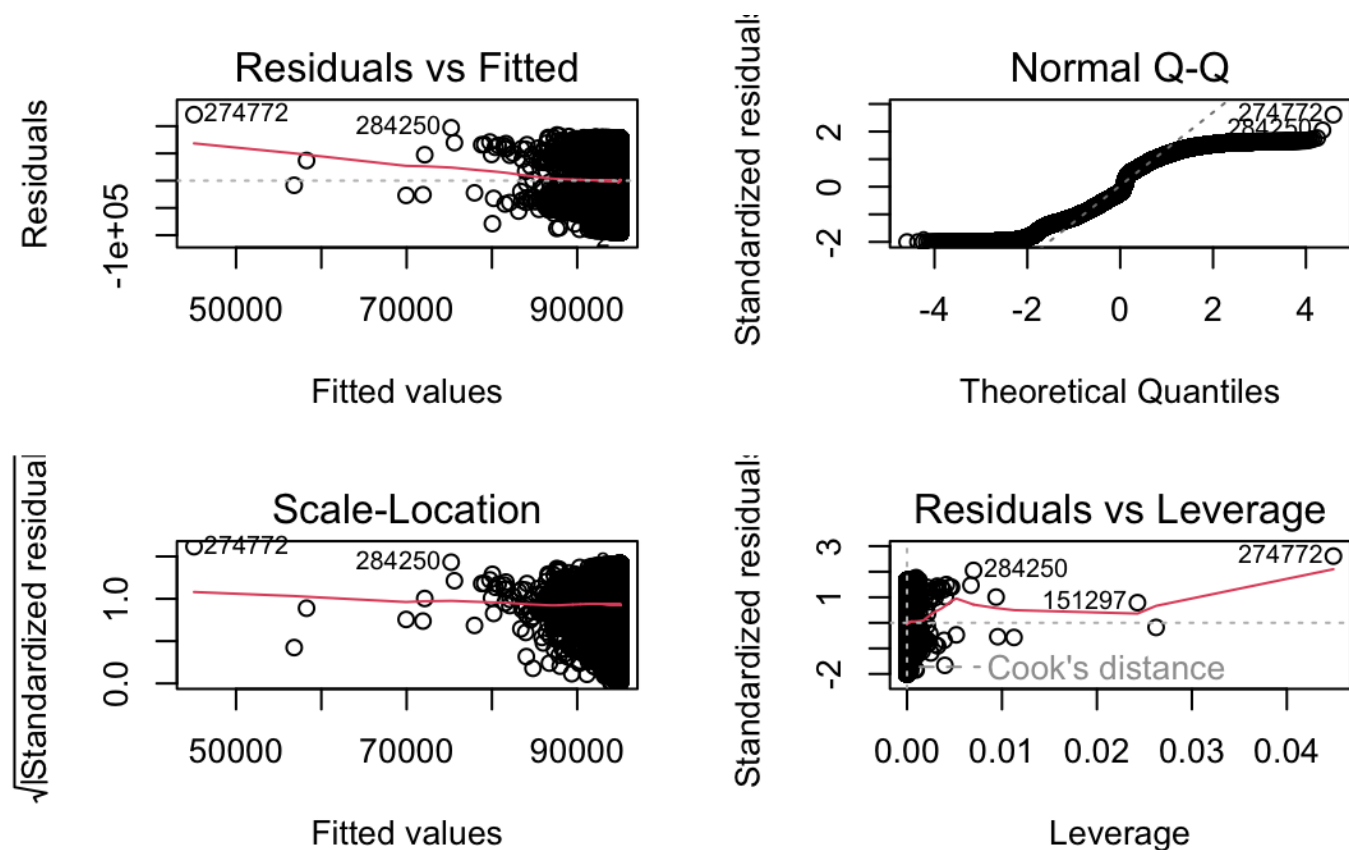
F-statistic: 24.52 on 1 and 227843 DF, p-value: 7.352e-07

From this summary we can learn that the p value is extremely small which represents that there is a strong relationship between time and amount. the multiple R squared value is very low which suggests that the model explains very little of the variance of time. the residual standard error is high which indicates that the prediction

would not be as accurate.

[Hide](#)

```
par(mfrow=c(2,2))
plot(lm1)
```


[Hide](#)

```
lm2<-lm(Time~Amount+Class, data = train)
summary(lm2)
```

Call:

```
lm(formula = Time ~ Amount + Class, data = train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-95081	-40570	-10082	44481	120752

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	9.509e+04	1.054e+02	901.939	< 2e-16 ***
Amount	-1.932e+00	3.929e-01	-4.918	8.77e-07 ***
Class	-1.374e+04	2.355e+03	-5.835	5.39e-09 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

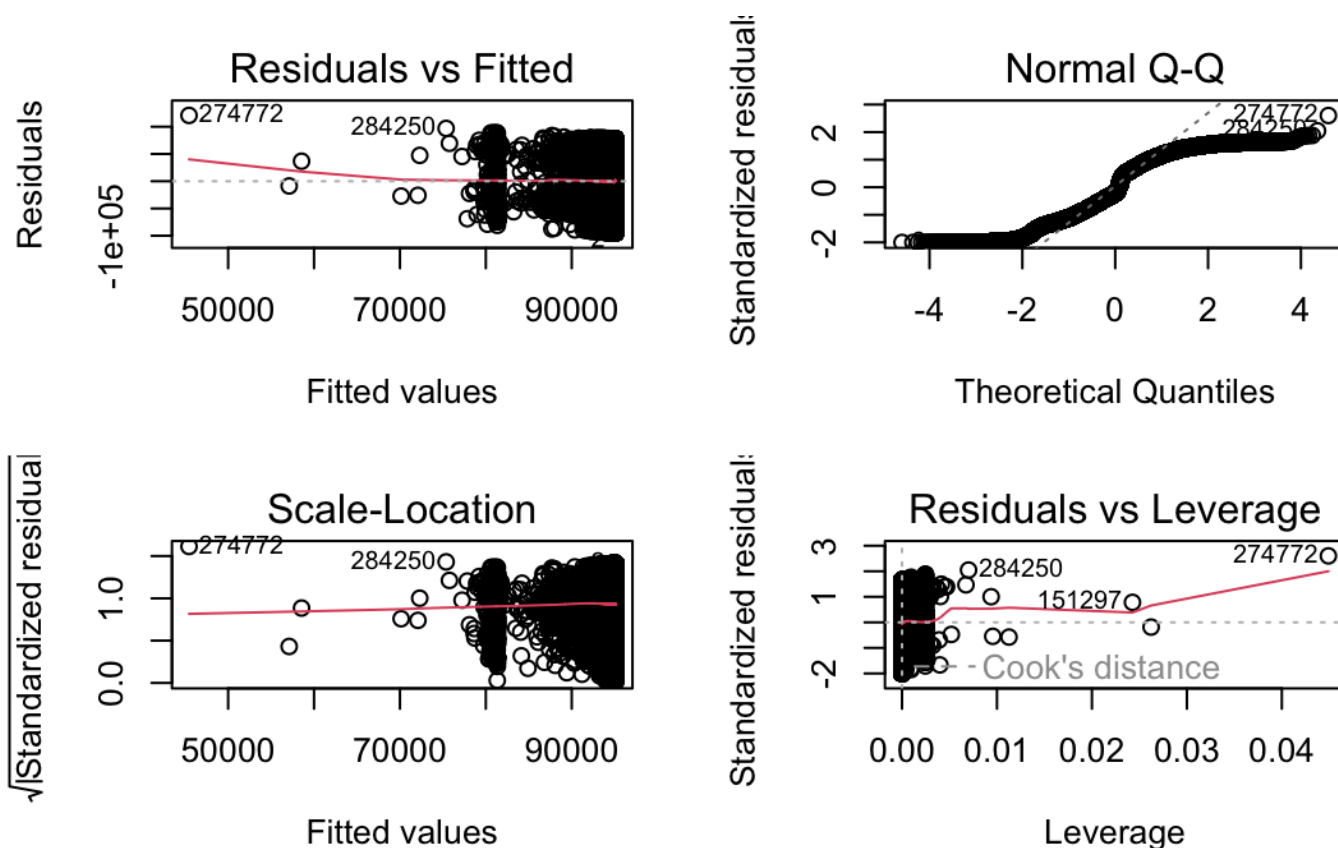
Residual standard error: 47470 on 227842 degrees of freedom

Multiple R-squared: 0.000257, Adjusted R-squared: 0.0002482

F-statistic: 29.29 on 2 and 227842 DF, p-value: 1.918e-13

Hide

```
par(mfrow=c(2,2))
plot(lm2)
```



Hide

```
lm3<-lm(Time ~ Time+Amount+Class, data = train)
```

Warning: the response appeared on the right-hand side and was dropped  
Warning: problem with term 1 in model.matrix: no columns are assigned

[Hide](#)

```
summary(lm3)
```

Call:

```
lm(formula = Time ~ Time + Amount + Class, data = train)
```

Residuals:

Min	1Q	Median	3Q	Max
-95081	-40570	-10082	44481	120752

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	9.509e+04	1.054e+02	901.939	< 2e-16 ***
Amount	-1.932e+00	3.929e-01	-4.918	8.77e-07 ***
Class	-1.374e+04	2.355e+03	-5.835	5.39e-09 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 47470 on 227842 degrees of freedom

Multiple R-squared: 0.000257, Adjusted R-squared: 0.0002482

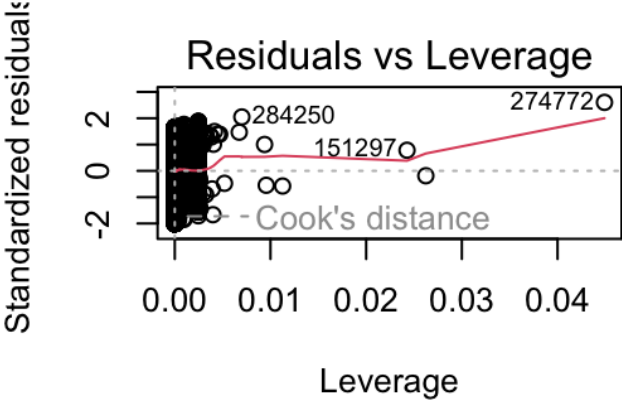



F-statistic: 29.29 on 2 and 227842 DF, p-value: 1.918e-13

[Hide](#)

```
par(mfrow=c(2,2))  
plot(lm3)
```

Warning: the response appeared on the right-hand side and was dropped  
Warning: problem with term 1 in model.matrix: no columns are assigned





Hide

```
pred1<-predict(lm1, newdata = test)
cor1 <- cor(pred1, test$Time)
mse1 <- mean ((pred1-test$Time)^1)
rmse1 <-sqrt(mse1)
print(paste('correlation: ',cor1))
```

```
[1] "correlation:  0.0115679423668671"
```

[Hide](#)

```
print(paste('mse: ',mse1))
```

```
[1] "mse:  385.643485196248"
```

[Hide](#)

```
print(paste('rmse: ',rmse1))
```

```
[1] "rmse:  19.6378075455548"
```

[Hide](#)

```
pred2<-predict(lm2, newdata = test)
cor2 <- cor(pred2, test$Time)
mse2 <- mean ((pred2-test$Time)^2)
rmse2 <-sqrt(mse2)
print(paste('correlation: ',cor2))
```

```
[1] "correlation:  0.0170486625902318"
```

[Hide](#)

```
print(paste('mse: ',mse2))
```

```
[1] "mse:  2257396538.92481"
```

[Hide](#)

```
print(paste('rmse: ',rmse2))
```

```
[1] "rmse:  47512.0672979488"
```

[Hide](#)

```
pred3<-predict(lm3, newdata = test)
```

Warning: prediction from a rank-deficient fit may be misleading

[Hide](#)

```
cor3 <- cor(pred3, test$Time)
mse3 <- mean ((pred3-test$Time)^2)
rmse3 <-sqrt(mse3)
print(paste('correlation: ',cor3))
```

```
[1] "correlation:  0.0170486625902318"
```

[Hide](#)

```
print(paste('mse: ',mse3))
```

```
[1] "mse:  2257396538.92481"
```

[Hide](#)

```
print(paste('rmse: ',rmse3))
```

```
[1] "rmse:  47512.0672979488"
```

Output: Fromt the correlation, mse and rmse we notice that the first linear regression model is the best one. We see that because of the low mse and rmse of the the first one compared to the other two. this may happen because the variables in the second and third models may not connect with the prediction. They divert the data in some other way.